

0. Review

Dong-Gyu, Lee

Dept. of Statistics, KU

2020, Apr 3

Contents

- 1 Analysis Goal
- 2 ILSVRC
- 3 FER2013
- 4 AlexNet
- 5 VGG
- 6 GoogLeNet
- 7 ResNet
- 8 Reference

Today's Goal

- Before studying the Jupiter code for AlexNet, VGG, GoogleNet, and ResNet, let's review the core of each model.

ILSVRC

- ImageNet Large Scale Visual Recognition Challenge(ILSVRC) was an annual computer vision contest held between 2010 and 2017.
- It's also called ImageNet Challenge.




Figure 1: ILSVRC(Classification)

- For this challenge, the training data is a subset of ImageNet: 1000 synsets, 1.2 million images.
- There are 50K images for validation and 150K images for testing.




Figure 2: ImageNet Dataset Used in the ILSVRC

- FER2013 dataset consists of 35,887 grayscale, 48x48 sized face images with various emotions - 7 emotions.
- Emotion labels in the dataset(imbalanced) : Angry, Disgust, Fear, Happy, Sad, Surprise, Neutral



Figure 3: ImageNet Dataset Used in the ILSVRC

- However, in FER2013 dataset, quite a few images are incorrectly labeled.
- So the FER Plus dataset was newly created in 2016.

1. AlexNet

- The number of parameters of AlexNet[4] is almost 62 million.







Figure 4: AlexNet

1. AlexNet

- Tanh vs ReLU.
- Dropout.
- Overlapping pooling.
- LRN.
- Data augmentation.



(a) ReLU



(b) Dropout

Figure 5: AlexNet Characteristic

2. VGG

- The number of parameters of VGG16[6] is almost 138 million.




Figure 6: VGG16 Structure

2. VGG

- Interested in deeper layers.
- Only 3x3 convolution.
- VGG11, VGG13, VGG16, VGG19.

ConvNet Configuration					
A	A-LRN	B	C	D	E
11 weight layers	11 weight layers	13 weight layers	16 weight layers	16 weight layers	19 weight layers
input (224 × 224 RGB image)					
conv3-64	conv3-64 LRN	conv3-64	conv3-64	conv3-64	conv3-64
maxpool					
conv3-128	conv3-128	conv3-128	conv3-128	conv3-128	conv3-128
maxpool					
conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256 conv1-256	conv3-256 conv3-256 conv3-256	conv3-256 conv3-256 conv3-256 conv3-256
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 conv1-512	conv3-512 conv3-512 conv3-512	conv3-512 conv3-512 conv3-512 conv3-512
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 conv1-512	conv3-512 conv3-512 conv3-512	conv3-512 conv3-512 conv3-512 conv3-512
maxpool					
FC-4096					
FC-4096					
FC-1000					
soft-max					

Figure 7: VGG

3. GoogLeNet

- The number of parameters of GoogLeNet[7] is almost 5 million.




Figure 8: GoogLeNet Structure

3. GoogLeNet

- 22 layers.
- Inspired by network in network(NIN)[5].
- Inception module(sparse & dense connectivity, 1x1 convolution).
- Auxiliary classifier(0.3:0.3:1).
- Global Average Pooling(GAP).

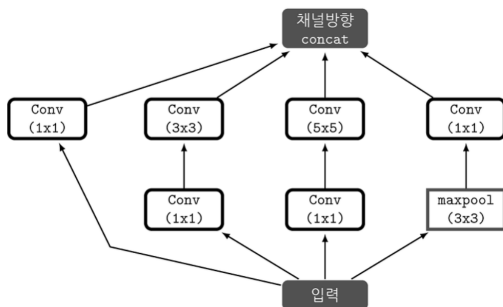




Figure 9: Inception Module

3. GoogLeNet

- Ref) Network in Network :



(a) Linear Convolution Layer



(b) Mlpconv Layer

3. GoogLeNet

- GoogLeNet Summary :

type	patch size/ stride	output size	depth	#1×1	#3×3 reduce	#3×3	#5×5 reduce	#5×5	pool proj	params	ops
convolution	7×7/2	112×112×64	1							2.7K	34M
max pool	3×3/2	56×56×64	0								
convolution	3×3/1	56×56×192	2		64	192				112K	360M
max pool	3×3/2	28×28×192	0								
inception (3a)		28×28×256	2	64	96	128	16	32	32	159K	128M
inception (3b)		28×28×480	2	128	128	192	32	96	64	380K	304M
max pool	3×3/2	14×14×480	0								
inception (4a)		14×14×512	2	192	96	208	16	48	64	364K	73M
inception (4b)		14×14×512	2	160	112	224	24	64	64	437K	88M
inception (4c)		14×14×512	2	128	128	256	24	64	64	463K	100M
inception (4d)		14×14×528	2	112	144	288	32	64	64	580K	119M
inception (4e)		14×14×832	2	256	160	320	32	128	128	840K	170M
max pool	3×3/2	7×7×832	0								
inception (5a)		7×7×832	2	256	160	320	32	128	128	1072K	54M
inception (5b)		7×7×1024	2	384	192	384	48	128	128	1388K	71M
avg pool	7×7/1	1×1×1024	0								
dropout (40%)		1×1×1024	0								
linear		1×1×1000	1							1000K	1M
softmax		1×1×1000	0								

Figure 11: GoogLeNet

4. ResNet

- 18,34,50,101,152 layers.
- The number of parameters of ResNet34[1],[2] is almost 21 million.






Figure 12: ResNet34 Structure

4. ResNet

- Degradation.
- Residual learning.
- Skip(Shortcut) connection.
- Bottleneck architecture.
- Batch Normalization[3].



(a) Training




(b) Test


Figure 13: Degradation Problem

4. ResNet

- Residual Learning :



(a) Simple Neural Network




(b) Skip Connection in ResNet


Figure 14: Residual Learning

4. ResNet

- Bottleneck Architecture :



(a) Block for ResNet-18/34



(b) Block for ResNet-50/101/152

Figure 15: Right: Bottleneck Architecture

4. ResNet


- ResNet Summary :

layer name	output size	18-layer	34-layer	50-layer	101-layer	152-layer
conv1	112×112	7×7, 64, stride 2				
conv2_x	56×56	3×3 max pool, stride 2				
		$\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$
conv3_x	28×28	$\begin{bmatrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 8$
conv4_x	14×14	$\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 23$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 36$
conv5_x	7×7	$\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$
	1×1	average pool, 1000-d fc, softmax				
FLOPs		1.8×10^9	3.6×10^9	3.8×10^9	7.6×10^9	11.3×10^9


Figure 16: ResNet

4. ResNet

- ResNet Result :



(a) Before



(b) After

Figure 17: ResNet Effect

- [1] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 770–778, 2016.
- [2] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Identity mappings in deep residual networks. In *European conference on computer vision*, pages 630–645. Springer, 2016.
- [3] Sergey Ioffe and Christian Szegedy. Batch normalization: Accelerating deep network training by reducing internal covariate shift. *arXiv preprint arXiv:1502.03167*, 2015.
- [4] Alex Krizhevsky, Ilya Sutskever, and Geoffrey E Hinton. Imagenet classification with deep convolutional neural networks. In *Advances in neural information processing systems*, pages 1097–1105, 2012.
- [5] Min Lin, Qiang Chen, and Shuicheng Yan. Network in network. *arXiv preprint arXiv:1312.4400*, 2013.

- [6] Karen Simonyan and Andrew Zisserman. Very deep convolutional networks for large-scale image recognition. *arXiv preprint arXiv:1409.1556*, 2014.
- [7] Christian Szegedy, Wei Liu, Yangqing Jia, Pierre Sermanet, Scott Reed, Dragomir Anguelov, Dumitru Erhan, Vincent Vanhoucke, and Andrew Rabinovich. Going deeper with convolutions. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 1–9, 2015.